

1. Develop Azure Compute Solutions

└ 1.1 Implement Containerized Solutions

└ 1.1.1 Create and Manage Container Images for Solutions

- What is the correct structure of a Dockerfile?
- How do you optimize a Dockerfile for size and performance?
- How do multi-stage builds work?
- What commands are used to build an image locally?
- How do you tag versions appropriately?
- What is the correct format for image names (registry/repository:tag)?
- How do tags work and how are they used in CI/CD pipelines?
- How do you create and configure an ACR?
- What are the SKU tiers and when do you use them?
- How do you push/pull images using Docker CLI?
- How do you authenticate to ACR (admin account, service principal, managed identity)?
- What is az acr login and when is it required?
- How does Azure App Service, Azure Kubernetes Service (AKS), and Azure Container Instances (ACI) consume images from ACR?
- What are the permission requirements for pulling from ACR?
- How do you scan images for vulnerabilities?
- What tools or services are used to harden container images?
- How do you manage image versions across environments?
- How do you clean up unreferenced or old images in ACR?
- How can ACR tasks automate image builds?
- How do you use az acr task to create scheduled or event-driven builds?
- What are the pros and cons of storing images publicly vs privately?
- How do you configure access control for image repositories?

What is the correct structure of a Dockerfile?

```
FROM <base_image>
[ENV <key>=<value> ...]
[WORKDIR <directory>]
[COPY <src> <dest>]
[RUN <command>]
[EXPOSE <port>]
[CMD ["executable", "param1", ...]]
```

Key Points:

- FROM must be first (defines base image).
 - WORKDIR sets working directory inside container.
 - COPY adds files to the image from build context.
 - RUN executes shell commands (e.g., install packages).
 - EXPOSE is optional metadata; doesn't actually open ports.
 - CMD defines default container startup command (only one allowed; last one wins).
-

How do you optimize a Dockerfile for size and performance?

1. **Use slim/minimal base images:**
Prefer FROM mcr.microsoft.com/dotnet/aspnet:7.0-alpine over full images.
3. **Leverage multi-stage builds:**
Build in one stage, copy only final output to runtime stage to reduce size.
4. **Minimize layers:**
Group related RUN commands and clean up temp files in the same layer:
4. **Avoid unnecessary files:**
Use .dockerignore to exclude files (e.g., .git, node_modules).
5. **Set only needed environment variables and permissions:**
Avoid excessive ENV or USER changes unless required.

How do multi-stage builds work?

Reduce final image size by separating build and runtime stages.

- Define multiple FROM statements in one Dockerfile.
- Use an alias for the build stage (AS build).
- Copy only needed artifacts from the build stage into the final image.

What commands are used to build an image locally?

`docker build -t <name>:<tag> <path>`

e.g. `docker build -t myapp:latest .`

How do you tag versions appropriately?

Use semantic tags like:

latest, v1.0.0, dev, staging

e.g. `docker build -t myapp:v1.0.0 .`

Tag meaning should reflect version or environment for clarity and traceability.

What is the correct format for image names?

`<registry>/<repository>:<tag> :`

`myregistry.azurecr.io/myapp:v1.0.0`

Registry is optional for local images. Tag defaults to **latest** if omitted.

How do tags work and how are they used in CI/CD pipelines?

Tags identify image versions. Pipelines use tags to pull, test, and deploy specific builds:

docker push myapp:staging → used in staging environment.

latest often used in dev, **versioned tags** in prod.

How do you create and configure an ACR?

- Create ACR: `az acr create --name <acr-name> --resource-group <rg> --sku Basic`
- Enable admin access: `az acr update -n <acr-name> --admin-enabled true`
- Login: `az acr login --name <acr-name>`

What are the SKU tiers and when do you use them?

- **Basic** – Dev/test, low-cost, limited features
 - **Standard** – Prod-ready, geo-replication support
 - **Premium** – High-scale, content trust, private endpoints, more throughput
-

How do you push/pull images from/to ACR using Docker CLI?

Push:

1. `docker tag myapp myacr.azurecr.io/myapp`
2. `docker push myacr.azurecr.io/myapp`

Pull:

`docker pull myacr.azurecr.io/myapp`

How do you authenticate to ACR (admin, service principal, managed identity)?

- **Admin account:**
Enable with `az acr update --admin-enabled true`, then use provided username/password.
 - **Service principal:**
Assign AcrPush/AcrPull role, login with docker login using SP credentials.
 - **Managed identity:**
Grant role to identity, Azure services (e.g., App Service) authenticate automatically.
-

What is az acr login and when is it required?

`az acr login --name <acr-name>`

Authenticates Docker CLI with ACR.

Required for manual Docker pushes/pulls. Not needed for Azure services using managed identity.

How does Azure App Service, Azure Kubernetes Service (AKS), and Azure Container Instances (ACI) consume images from ACR?

- **App Service:** Configure container settings with ACR URL; use managed identity or admin acc.
 - **AKS:** Enable ACR integration via `az aks update` or use `imagePullSecrets`.
 - **ACI:** Reference image with full ACR path; grant access via identity or admin credentials.
-

What are the permission requirements for pulling from ACR?

The identity must have **AcrPull** role on the ACR.

Can be assigned to:

- User
 - Service principal
 - Managed identity (App Service, AKS, etc.)
-

How do you scan images for vulnerabilities?

Use **Microsoft Defender for Cloud** with ACR integration. It scans images on push and shows CVEs in the portal. Enable under *Defender plans > Container registries*.

What tools or services are used to harden container images?

- **Microsoft Defender for Cloud** – vulnerability scanning
 - **Dockerfile best practices** – minimize layers, use minimal base images
 - **Content trust** – ensure image integrity
 - **Private ACR** – restrict access
 - **ACR Tasks** – automate secure builds
-

How do you manage image versions across environments?

Use consistent **tagging strategy** (e.g., dev, staging, v1.0.0).

Promote images by **re-tagging** and pushing to ACR for each stage.

How do you clean up unreferenced or old images in ACR?

Use **ACR Tasks with retention policies** or manual cleanup via:

az acr repository delete --name <acr> --image <repo>:<tag>

How can ACR tasks automate image builds?

ACR Tasks can auto-build images on source code or base image changes:

az acr task create with **--source** and **--cmd "docker build"**

Supports triggers (e.g., Git push) and scheduling.

How do you use az acr task to create scheduled or event-driven builds?

Event-driven:

```
az acr task create \  
  --name mytask \  
  --registry myacr \  
  --image myapp:{{.Run.ID}} \  
  --context https://github.com/org/repo.git \  
  --file Dockerfile \  
  --git-access-token <token>
```

Scheduled:

```
az acr task create \  
  --name mytask \  
  --registry myacr \  
  --schedule "0 2 * * *" \  
  --image myapp:nightly \  
  --context https://github.com/org/repo.git \  
  --file Dockerfile
```

What are the pros and cons of storing images publicly vs privately?

Public: Easy access, no auth needed — but insecure, no access control.

Private: Secure, controlled access — but needs auth, may cost more.

How do you configure access control for image repositories?

Assign **AcrPull** or **AcrPush** roles to users, service principals, or managed identities using **Azure RBAC** on the ACR resource.